

Introduction to parallel and distributed computing in high-level programming languages: **MATLAB** and **Julia**

Henryk Modzelewski (hmodzelewski@eoas.ubc.ca)

&

Keegan Lensink (klensink@eoas.ubc.ca)

EOAS SLIM/WFRT teams

WestGrid Research Computing Summer School - UBC - June 2017

Hopeful agenda

- Some preparation - WestGrid access and setup
- Why high-level languages?
- Why Parallel & Distributed?
- Basics: MATLAB vs. Julia
- Parallel/Distributed features: MATLAB vs. Julia
- Debugging and profiling
- Extensions via other languages: MATLAB vs. Julia
- Some simple examples and exercises

WestGrid access & setup

- Host: `$ ssh -Y user_name@grex.westgrid.ca`
- Get resources:
 - `$ qsub -l -l nodes=1:ppn=3;walltime=3:00:00`
- Environment modules:
 - `$ module load python/2.7.9-gcc52`
 - `$ module load julia/0.5.2-bin`
- `$ export JULIA_LOAD_PATH=.`

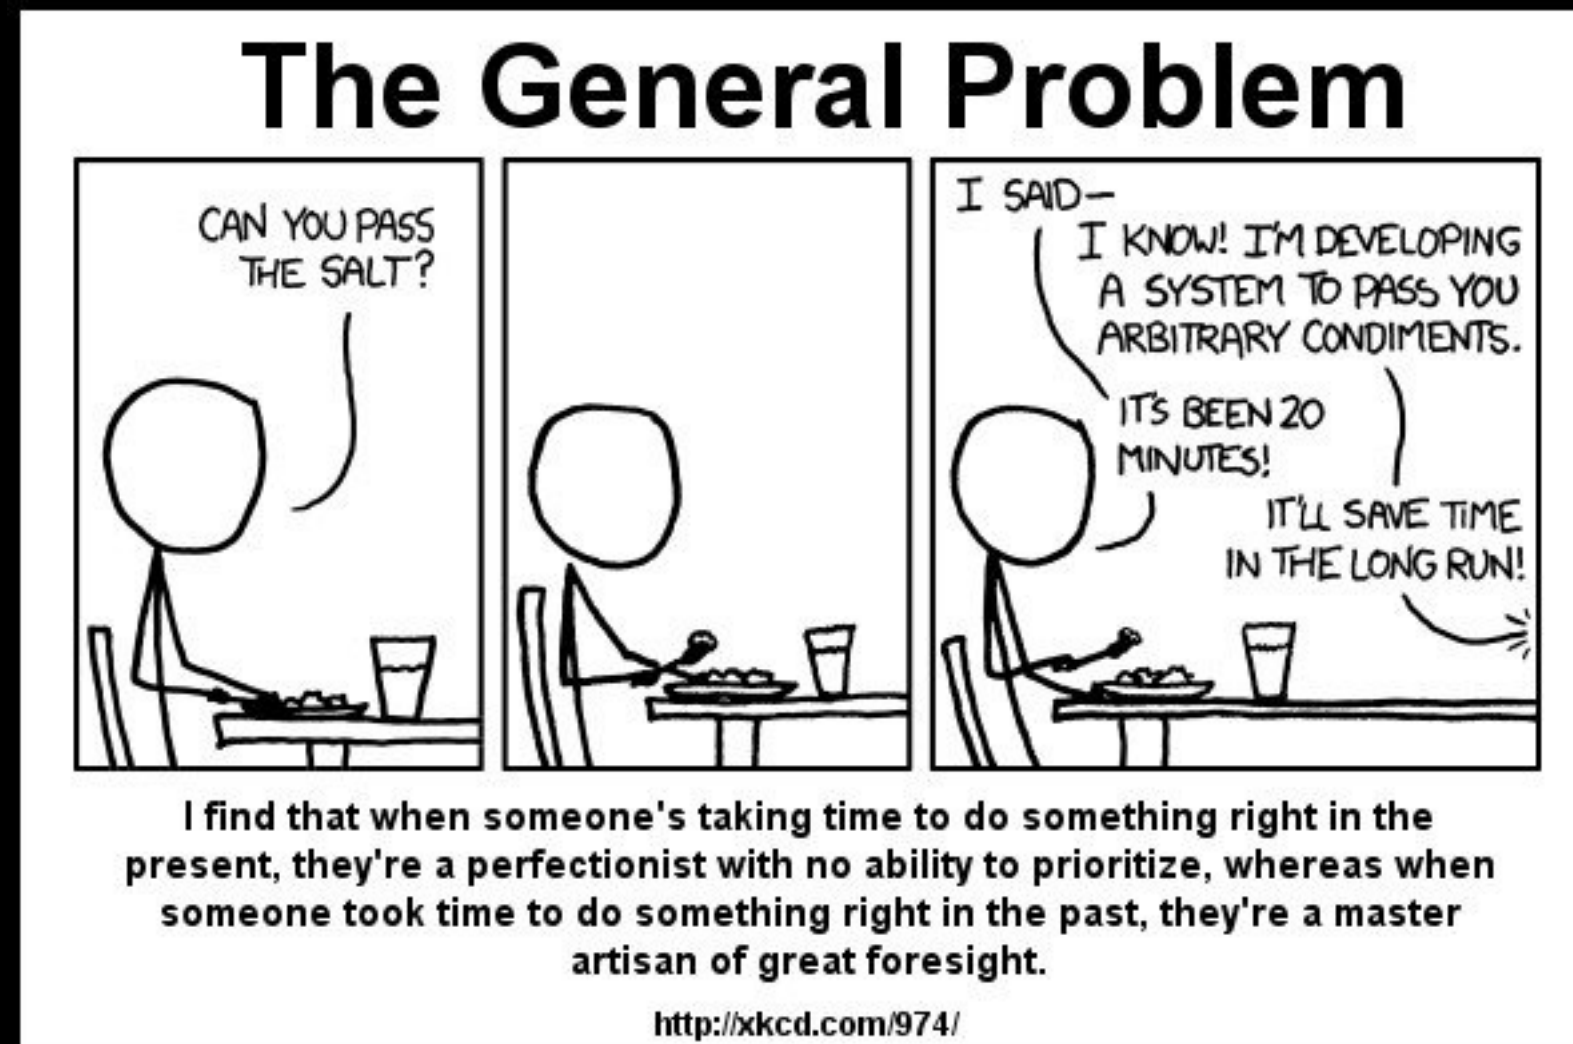
WestGrid access & setup

- Clone GIT repo
 - <https://github.com/henryk-modzelewski/MvsJ.git>
- Add packages (in \$ julia)
 - `Pkg.add("BenchmarkTools")`
 - `Pkg.add("Gallium")`
 - `Plg.add("DistributedArrays")`

Why high-level languages?

Why high-level languages?

- To develop faster
- To build less code
 - utilize libraries build by others
- To debug and profile easier
- Natively domain-specific languages (BLA)
 - or easy to adapt (OOP)
 - or having supporting packages
- To run slower :) or maybe not.



Pareto principle
80/20 (?)

Propaganda!!





Figure: benchmark times relative to C (smaller is better, C performance = 1.0).

C and Fortran compiled with gcc 5.1.1. C timing is the best timing from all optimization levels (-O0 through -O3). C, Fortran and Julia use [OpenBLAS](#) v0.2.14. The Python implementations of rand_mat_stat and rand_mat_mul use NumPy (v1.9.2) functions; the rest are pure Python implementations. Plot created with [Gadfly](#) and [IJulia](#) from [this notebook](#).

Why parallel and distributed?

- To run applications faster
- To process more data
- To run bigger problems
- **Above subject to potential for scaling**
- Serial vs parallel portion of the code
- Miracles do not happen
 - they require hard work
 - or just do not happen :)



Basics - Big picture

MATLAB vs Julia - big picture

	MATLAB	Julia
License	Proprietary	Open Source
Cost	License	Free
Libraries	Core licensed + free	Core free + licensed
Documentation	In application & on-line	In application & on-line
Installation	Requires admin	User installable

MATLAB vs Julia - big picture

	MATLAB	Julia
Robustness	Old foe minor changes	New friend changes quite a lot
Scalability	Supported up to 256*	Supported yet unknown limits
Runtime	Interpreter	JiT compiler
OOP	Supported via classes single-dispatch	Supported via types multiple-dispatch
Garbage collection	Automatic	Automatic

MATLAB vs Julia - big picture

	MATLAB	Julia
Profiling	Easy and thorough	Basic profiling tools
Debugging	Easy and flexible GUI support	Still quite crude CLI only

Resources

- MATLAB

- <https://www.mathworks.com/help/matlab/>
- <http://www.mathworks.com/matlabcentral/>
- <https://www.youtube.com/user/MATLAB>

- Julia

- <https://julialang.org>
- https://en.wikibooks.org/wiki/Introducing_Julia
- <https://www.youtube.com/user/JuliaLanguage>
- <https://github.com/JuliaLang>

Basics - running front-end

Basics - running

- In workspace
 - use command:
 - `$ matlab [-nodesktop | -display]`
 - execute from there
 - type function name for external script
 - finish with exit
- From shell
 - use command:
 - `$ matlab [-nodesktop | -display] -r 'script/function_name[; exit]'`

Basics - running

- In workspace
 - use command:
 - `$ julia`
 - execute from there
 - use `include()` for script
 - finish with `^D` or `quit()`
- From shell
 - use command:
 - `$ julia script_name`
 - `$ julia -L lib_name scrip_name`
 - shebang and execute perms
 - `#!/usr/bin/env julia`
 - `$ chmod u+x script_name`

Using workspace

	MATLAB	Julia (REPL)
Line suppression	...;	...;
SHELL escape	!	;
Help	help	?
Search		apropos
Locate	which	
Reset	clear [...]	workspace()

Basics - language fundamentals

Basics - assignments and operators

- Both languages support the typical syntax used in other languages
- Assignments look the same
- Similar operators
- Standard operator precedence
 - same precedent asses left-to-right
- Function calls look the same

```
a=1
```

```
b=a+1
```

```
c=a*b
```

```
x=foo(y,a,b,c)
```

Build-in simple types

- Numerals:
 - integers with different precision
 - floats/complex with single/double precision
 - hard to predict/control the outcome
 - computations are cared in single or double precision (!)
 - typically downgrades to single in binary operators (!)
- Strings (characters are size (1,1) strings)
- Date & Time
- ...

Build-in simple types

- Numerals:
 - integers with different precision
 - floats/complex with different precision
 - infers result of operation from higher-precision components
 - `promote_type()`
- Strings
- Characters (null size)
- Date & Time
- ...

Built-in complex types

- Numerical arrays column-major, indexed (i,...)
 - (either double or single)
 - support slicing with ':' e.g. (:,i:j)
- Cell arrays (single dim) - hold anything, indexed {i}
 - slicing {i:j} returns separate elements
- Structures - expandable
- Composite() - for parallel computations

Built-in complex types

- Numerical arrays column-major, indexed `[i,...]`
 - Matrix and Vector are different types of Array
- “Any” arrays - anything
- Tuples, immutable, indexed with `[i]`
 - above support slicing with `:`, e.g. `[i:j,:]`

Basics - some notable differences

Basics - some notable differences

- numeral assignments always create type 'double'
- complex using 'j'
- scalars/characters have size (1,1)
- boolean (true/false) negate with '~'
- strings surrounded by 'double quote'
- characters surrounded by 'single quote' (!)
- strings are arrays of characters
- string concatenation via function strcat
- comments start with '%'

```
a=1  
whos a  
b=1.  
whos b
```

```
d=1+j*4
```

```
size(a)
```

```
T=true  
F=~T
```

```
s='hello'  
c=s(1)
```

```
ab=strcat('a','b')  
ab=['a','b']
```

Basics - some notable differences

- numeral assignments default integer/float/complex types, other type require `convert()`
- complex using 'im'
- scalars/characters have null size
- boolean (true/false) negate with '!'
 - strings surrounded by "double quote"
 - characters surrounded by 'single quote'
 - strings are arrays of characters
 - strings concatenate via '*'
 - strings can be interpolated
 - in-place functions - names end with '!' convention
 - `+=`, `-=`, `*=`, `/=`
 - comments start with '#'

```
a=1
typeof(a)
b=1.
typeof(b)
```

```
d=1+im*4
```

```
size(a)
```

```
T=true
F=!true
```

```
S="test"
C='c'
```

```
S2="test | "*"test2"
S3="This is $S2"
S4="This is $a"
```

```
A=[1, 2]
push!(A,3)
```

Basics - flow control

Basics - flow control

- for loops
 - for $i=1:10$; ...; end
 - for $i=1:2:13$; ...; end
 - continue & break
- while $cond$; ...; end
 - continue
- switch EXP ; case exp ...; otherwise ...; end
 - no break
- if $cond$... ; elseif $cond2$...; else... ; end
- exceptions
 - throw() & rethrow()
 - try ...; catch ...; end

Basics - flow control

- for loops
 - for $i=1:10$
 - for $i=1:2:11$
 - for e in 'iterable'
 - continue & break
- while $cond; \dots; end$
 - continue & break
- if $cond \dots; elseif\ cond2 \dots; else \dots; end$
 - if $cond \dots end$
 - ternary operator: $x = cond ? \dots : \dots$
- exceptions
 - throw() & rethrow()
 - try $\dots; catch; end$

Basics - function

Functions

- One accessible function per file
- Arguments
 - required (positional)
 - optional (positional)
 - keywords 'key','value' sequence
- No type matching
 - processing arguments in order
 - or use built-in parser
- In-place functions
- Anonymous functions
- Function pointers with '@'

```
% my function file
```

```
%main
```

```
function [x,y]=foo(a,b,c,'key','value')
```

```
...
```

```
end
```

```
%private function
```

```
function c=extra(d,e)
```

```
...
```

```
end
```

```
% private in-place
```

```
function c=inplace(c)
```

```
...
```

```
end
```

```

function op = opCurvelet(m, n, nbscales, nbangles,...
    finest,ttype, is_real)

    assert( isscalar(m) && isscalar(n),['Please ensure'...
        ' sizes are scalar values']);
    if nargin < 3, nbscales = max(1,ceil(log2(min(m,n)) - 3)); end;
    if nargin < 4, nbangles = 16;                end;
    if nargin < 5, finest = 0;                    end;
    if nargin < 6, ttype = 'WRAP';                end;
    if nargin < 7, is_real = 1;                    end;
    assert( strcmp(ttype,'WRAP') || strcmp(ttype,'ME'),...
        ['Please ensure ttype is set correctly. Options are'...
        ' "WRAP" for a wrapping transform and "ME" for a'...
        ' mirror-extended transform']);
    assert( isscalar(nbscales) && isscalar(nbangles),...
        'Please ensure nbscales and nbangles are scalar values');
    assert((any(finest == [0 1 2])) && (is_real==0||is_real==1),...
        'Please ensure finest and is_real are appropriate values');
    if finest==0, assert( nbscales>1, ['Please ensure that '...
        'm and n are large enough for nbscales to be '...
        'greater than 1 while finest is set to 0']);
    end
...

```

```

function x = poMatCon(pathname,varargin) % Constructor for poMatCon

    % Parse param-value pairs using input parser
    p = inputParser;
    p.addParamValue('precision','double',@ischar);
    p.addParamValue('repeat',0,@isscalar);
    p.addParamValue('readonly',0,@isscalar);
    p.addParamValue('distribute',0,@isscalar);
    p.addParamValue('copy',0,@isscalar);
    p.KeepUnmatched = true;
    p.parse(varargin{:});

    if (isdir(pathname)) % Loading file
        if (p.Results.copy == 0) % overwrite case
            headerIn = SDCpkg.Reg.io.NativeBin.serial.HeaderRead(pathname);
            td = pathname;
        else % no overwrite
            td = SDCpkg.Reg.io.makeDir();
            SDCpkg.Reg.io.NativeBin.serial.FileCopy(pathname,td);
            headerIn = SDCpkg.Reg.io.NativeBin.serial.HeaderRead(td);
        end
    else
        ...
    end

```


Functions

- No limit on # of functions in file
- From weak to strong typing
- Arguments
 - required (positional)
 - optional (positional until done)
 - keywords 'key'='value' after ';'
 - type matching and all above
 - still a bit shaky
- Returns last-line result or via return statement
- Anonymous functions
- Functions are objects

```
function foo(a,b,c=1,d=7;verb=true)
```

```
...
```

```
  return res
```

```
end
```

```
function extra(d::Int,e::Float64;verb::Bool=true)
```

```
...
```

```
  res = something
```

```
end
```

```
function extra{DT}(d::DT,e::DT)::DT
```

```
...
```

```
  res
```

```
end
```

```

function joCoreBlock(ops::joAbstractLinearOperator...;kwargs...)
    isempty(ops) && throw(joCoreBlockException("empty argument list"))
    l=length(ops)
    for i=1:l
        deltype(ops[i])==deltype(ops[1]) || throw(joCoreBlockException("domain type mismatch for $i operator"))
        reltype(ops[i])==reltype(ops[1]) || throw(joCoreBlockException("range type mismatch for $i operator"))
    end
    mykws=Dict(kwargs[i][1]>=>kwargs[i][2] for i in 1:length(kwargs))
    mo=Base.deepcopy(get(mykws, :moffsets, zeros{Int,0}))
    typeof(mo)<:AbstractVector || throw(joCoreBlockException("moffsets must be a vector"))
    eltype(mo)<:Integer || throw(joCoreBlockException("moffsets vector must have integer elements"))
    (length(mo)==l || length(mo)==0) || throw(joCoreBlockException("length of moffsets vector does not match number of operators"))
    no=Base.deepcopy(get(mykws, :noffsets, zeros{Int,0}))
    typeof(no)<:AbstractVector || throw(joCoreBlockException("noffsets must be a vector"))
    eltype(no)<:Integer || throw(joCoreBlockException("noffsets vector must have integer elements"))
    (length(no)==l || length(no)==0) || throw(joCoreBlockException("length of noffset vector does not match number of operators"))
    ws=Base.deepcopy(get(mykws, :weights, zeros{0}))
    typeof(ws)<:AbstractVector || throw(joCoreBlockException("weights must be a vector"))
    (length(ws)==l || length(ws)==0) || throw(joCoreBlockException("length of weights vector does not match number of operators"))
    name=get(mykws, :name, "joCoreBlock")
    typeof(name)<:String || throw(joCoreBlockException("name must be a string"))
    ME=get(mykws, :ME, 0)
    typeof(ME)<:Integer || throw(joCoreBlockException("ME must be Integer"))
    ME>=0 || throw(joCoreBlockException("ME must be >=0"))
    NE=get(mykws, :NE, 0)
    typeof(NE)<:Integer || throw(joCoreBlockException("NE must be Integer"))
    NE>=0 || throw(joCoreBlockException("NE must be >=0"))
    ...

```

Functions

- From weak to strong typing
- Parametric types
- Type stability
 - ensuring that compiler has least work to guess types
 - types assessed on the function boundary
- **Type stability is paramount for performance**
 - `@code_warntype`

```
function jo_convert{VT<:Integer}(DT::DataType,vin::AbstractArray{VT},warning::Bool=true)
```

```
    DT==VT && return vin
    if DT<:Integer
        if typemax(DT)>typemax(VT)
            vout=convert(AbstractArray{DT},vin)
        else
            throw(joUtilsException("jo_convert: Refused conversion from $VT to $DT."))
        end
    else
        vout=convert(AbstractArray{DT},vin)
    end
    return vout
end
```

```
function jo_convert{VT<:AbstractFloat}(DT::DataType,vin::AbstractArray{VT},warning::Bool=true)
```

```
    DT==VT && return vin
    if !(DT<:Integer)
        vout=convert(AbstractArray{DT},vin)
    else
        throw(joUtilsException("jo_convert: Refused conversion from $VT to $DT."))
    end
    return vout
end
```

```
function jo_convert{VT<:Complex}(DT::DataType,vin::AbstractArray{VT},warning::Bool=true)
```

```
    DT==VT && return vin
    if DT<:Complex
        vout=convert(AbstractArray{DT},vin)
    elseif DT<:AbstractFloat
        (warning && jo_convert_warn) && warn("jo_convert: Inexact conversion from $VT to $DT. Dropping imaginary part.")
        vout=convert(AbstractArray{DT},real(vin))
    else
        throw(joUtilsException("jo_convert: Refused conversion from $VT to $DT."))
    end
    return vout
end
```

Demonstration: in directory Loop

Exercise: in directory Clip

Use clip.m to produce Julia version in myclip.jl
- watch for type stability

Basics - OOP

Basics - OOP

- Inheritance, encapsulation, and polymorphism
- Classes
 - properties (public, protected, and private)
 - methods (public, protected, and private)
 - value and handle classes
- Single-dispatch - methods bound to class object

Basics - OOP

- Subtyping (type tree and associations) and ad hoc polymorphism (function overloading)
- Types
 - mutable or immutable
 - constructors (outer and inner)
 - abstract types
 - type aliases
- Methods (functions)
- Multiple-dispatch - methods bound by function's signature

Basics - OOP

- Modules
 - full-protection of module types and variables
 - semi-protection of methods (overloading for different types)

Real code (domain-specific language):

<https://github.com/slinggroup/JOLI.jl.git>

Demonstration: in directory OOP

Exercise: in directory OOP

Use figures.jl

1. addSquare

2. add >, >=, <=, ==

Adding and extending code

Adding 3-rd party code

- No package manager
 - Typically distributed as tar or zip archives
 - Installs anywhere
 - Just unzip/un-tar the archive somewhere and
 - add location to the path (next slides)
- Typical places:
 - ~/matlab
 - ~/Documents/MATLAB

Adding 3-rd party code

- Package Manager
 - GIT repositories
 - Installs into `~/.julia/v#.#/pckg_name`
 - `using` & `import`
- zip/tar archives
 - Just unzip/un-tar the archive somewhere and
 - add location to the path (next slides)

Adding 3-rd party code

- Registered packages
 - `Pkg.add("pckg_name")`
- Unregistered packages from GIT repositories
 - `Pkg.clone("pckg_url"[,local_name])`
- Checking out branches
 - `Pkg.checkout("pckg_name","branch_name")`

Organizing your source code

- Functions
- Classes
- Packages
 - SHELL environment MATLABPATH
 - `addpath("/home/me/Matlab"['-begin' '-end'])`
 - `addpath /home/me/Matlab [-begin|-end]`
- **MATLAB will find everything on the path**

Organizing your source code - Packages

Directory listing

- +SLIM_APPS
- +SLIM_APPS/available.m
- +SLIM_APPS/+tools
- +SLIM_APPS/+tools/Miscellaneous.m
- +SLIM_APPS/+tools/+algorithms
- +SLIM_APPS/+tools/+algorithms/REPSI.m
- +SLIM_APPS/+tools/+algorithms/TimeModeling.m
- +SLIM_APPS/+tools/+utilities
- +SLIM_APPS/+tools/+utilities/pSPOT.m
- +SLIM_APPS/+tools/+utilities/SPOT_SLIM.m

Packages create namespaces

```
>> SLIM_APPS.available
```

```
>> SLIM_APPS.tools.Miscellaneous(...)
```

```
>> SLIM_APPS.tools.utilities.pSPOT(...)
```

Organizing your source code

- Modules
 - SHELL environment `JULIA_LOAD_PATH`
 - append: `push!(LOAD_PATH, "/home/me/julia")`
 - prepend: `unshift!(LOAD_PATH, "/home/me/julia")`
 - **using & import will find all modules on the path**
- Functions (not recommended)
 - `include()` & `require()` & `reload()`
 - include mostly for organizing modules (relative or full path)
 - **require() & reload() will find functions on the path**

Parallel/Distributed features

Scalability - vertical vs. horizontal scaling

- vertical
 - expanding resources on single computing server
 - shared memory
- horizontal
 - adding more computing servers
 - distributing memory

Scalability - strong vs. weak scaling

- strong
 - expanding resources for same problem size
 - decreasing problem size per process
 - $(\text{sub-size} = \text{size} / \text{workers})$
- weak
 - increase problem proportional to increased resource
 - constant problem size per process
 - $(\text{size} = \text{sub-size} * \text{workers})$

Parallel tools

- No explicit multi-threading
- Communication
 - Master-worker model from main client
 - MPI mode between workers
- parfor loop (no nesting)
- spmd block (no nesting)
- distributed/codistributed arrays
 - distributed array constrictors
 - codistributed array builders

Parallel tools

- Functions:
 - parpool (interactive)
 - batch (non-interactive)
 - script submission
 - function call

Parallel tools

- Parallel invocation via functions:
 - `parpool(#)` - interactive
 - `batch` - noninteractive jobs
 - can also execute function remotely

Parallel tools

- Multi-threading
 - Implicit for some native components
 - Threads module (experimental)
 - `Threads.@threads`
- Communication between workers - RPC
 - “Master”-workers model (actor?)
- Global parallel calls
 - `@parallel` for (SharedArrays)
 - `pmap`

Parallel tools

- Feature - response to
 - `@spawn` & `@spawnat`
 - `remotecall()` - non-blocking
 - `remotecall_wait()` - blocking
 - `remotecall_fetch()` - blocking
 - `wait()` & `fetch()`

Parallel tools

- RemoteChannel - rewritable communication path
- @sync & @async for more sophisticated control
- Distributed arrays - Package DistributedArrays
 - spmd functionality in next v0.6 version of Julia

Parallel tools

- Thread control:
 - set threads: `$ export JULIA_NUM_THREADS=#`
 - BLA operators: `> BLAS.set_num_threads(#)`
- Parallel invocation:
 - command line: `$ julia -p # [--machinefile <file>]`
 - `addprocs()` and `rmprocs()`

Example; parallel command execution
in CodeExamples

Debugging and profiling - Keegan Lensink

Extensions via other languages: MATLAB vs. Julia

- MATLAB coder - compiling native MATLAB code
- Java - MATLAB is partly Java, and supports Java classes
- C/Fortran API - a bit cumbersome to use
 - mex extensions - great for
 - accelerating
 - OpenMP/multithreading

Extensions via other languages: MATLAB vs. Julia

- Python - Package PyCall
- MATLAB - Package MATLAB
- C/Fortran via ccall (dlopen)
 - great for
 - accelerating
 - OpenMP/multithreading

Advection examples - code analysis

	MATLAB	Julia
Serial	advection.m	advection.jl
Threaded		advectionT.jl
parfor @parallel	advectionPF.m advectionPFV.m	advectionPS.jl advectionPSS.jl
pmap		advectionPMAP.jl
spmd	advectionD.m	

The End